Alexander Sorrell

AI Principal / Builder / Breaker — Verification & Evaluation · (817) 996-6123 · alexander.sorrell.it@gmail.com

linkedin.com/in/alexander-sorrell-it · github.com/Alexander-Sorrell-IT

Colleyville, TX (Dallas–Fort Worth) — open to remote, onsite, or relocation

github.com/Alexander-Sorrell-IT/starreckon
npx starreckon

VERIFIED AUDIT · 30.1
MASTERWORK

PROFESSIONAL SUMMARY

Models propose. A deterministic function owns the verdict. Every decision ships with a receipt a stranger can re-check — a rule ID, a cryptographic hash chain, or the OS kernel itself.

PROOF OF WORK

- **Calibrated abstention under attack.** I ran live prompt injection against my own forensics court to find out what it does when someone is actively trying to poison the judge. F1 halved. Planted false positives stayed at zero, and every degraded run abstained. It fails silent, not wrong.
- **Two counters, two authors, one answer.** My token counter (starforge) and an independently written one by a different author — different language, no shared code — read the same 20,217 transcripts from 5 computers, each in its own container with the corpus read-only. They agree exactly on sessions, active days, hours, input and cache tokens, and differ by 0.07% on output. The designed failure mode is INCONCLUSIVE, never PASS.
- **The kernel certifies the boundary, not the code.** A socket probe executes inside the sandbox and outside it; the OS returns EPERM rather than relying on an application-level boolean. Reproducible at your desk — npx starreckon prove prints the exact command.
- **Byte-for-byte ground truth.** Clean-room reconstruction of an NVIDIA HGX factory tool's output to an identical md5 against the closed-source Windows original, non-circular by control. Replaying 10 real H100 boards through the rebuilt suite matches the factory's disposition 10/10, and it caught two boards shipped as BBX Pass while carrying 8 row-remaps.
- **The payout function cannot receive an AI verdict.** Re-derives policy deterministically server-side so nothing upstream can trick it into an unauthorized state. 8 of 8 scripted adversarial attacks held, zero violations.
- **Measured, not remembered.** 5 public repositories since August 2026: 13,896 lines, 174 named tests, 73 commits on 19 active days — counted by my own repo auditor, which strips vendored libraries, other people's code and duplicate copies before it counts anything.

ARTIFICIAL INTELLIGENCE & TECHNICAL SKILLS

AI / ML & Eval Science	LLM-as-judge pipelines; golden-set design & contamination control; self-preference bias calibration; principled abstention; adversarial evaluation & red-teaming; multi-agent orchestration; deterministic QA; MCP
Languages	Python, TypeScript / JavaScript, Rust (→ WASM), Solidity, Bash, C
Platforms & Systems	Splunk / SPL, Foundry, Docker, GitHub Actions CI, Linux namespaces (unshare) / sandboxing, PyPI & npm publishing, MCP servers, Next.js / React, Playwright, REST APIs, SDK design
Infrastructure & Security	NVIDIA DGX/HGX A100/H100 bring-up, diagnostics & firmware (retimer, NVSwitch, FPGA registers), Linux, macOS sandbox confinement, smart-contract & web vulnerability analysis, DFIR

PROFESSIONAL EXPERIENCE

Founder & Principal AI Systems Architect — SignalCore Inc (DE) & Bastion Fold LLC (NM) · Remote · 2025 — Present

Verification infrastructure for AI systems: deterministic evaluation harnesses, agent guardrails, and audit trails where the component that decides is physically separate from the component that acts. Multi-agent orchestration, LLM-as-judge pipelines, MCP servers written from scratch, adversarial red-teaming. SignalCore Inc (DE, Jul 2025) builds agent governance; Bastion Fold LLC (NM, Apr 2025) operates deterministic security verification. Shipped 11 packages and 5 public repositories of my own on github.com/matrixbuilderops (73 commits on 19 active days); won the T3 ADK Bounty Challenge Launch Edition with On the Record. Python, TypeScript, Rust, Solidity.

Frontier-Model Red Team & Evaluation — Handshake AI (Project Vox), Mercor (Les Artistes) & Alignerr · Remote · 2026 contracts

Adversarial prompt red-teaming against frontier models, rubric-based response assessment, reference-answer authoring, and multi-turn reasoning data curation used to harden model alignment and truthfulness. Authored paid ML-environment tasks for Alignerr's ML Envs pipeline.

Server Repair Technician L1 → L2 — OnTrack Staffing, then SMS Infocomm · Jan 2024 — Apr 2025

Ran the full factory sort flow on NVIDIA HGX/DGX A100 and H100 AI-compute systems — system check, BBX, functional stress, in-system run-in — isolating GPU-vs-baseboard faults with golden-part swaps and flashing re-timer and NVSwitch firmware; trained incoming technicians. A sort line is a machine for refusing to believe hardware is good until an independent test agrees. That is the exact discipline I build into evaluation software.

Junior Claims Evaluator — Claimplus · Apr 2021 — Jan 2024

Reviewed and validated property-insurance claims data; structured evaluation workflows and deterministic policy checks.

BY THE NUMBERS (MEASURED ACROSS EVERY PUBLIC REPOSITORY)

- **5 public repositories** of my own on github.com/matrixbuilderops: 13,896 lines, 174 named tests, 73 commits on 19 active days, 2026-08-05 → 2026-10-31.
- **By language:** Python 7,428, JavaScript 3,510, TypeScript 2,958 lines.
- *How it is counted:* bundled libraries, other people's code, generated files and duplicate copies are excluded (0 lines left out); a file in more than one repo counts once; commits count only mine and my agents'.

PROJECTS — EVERY REPOSITORY, WRITTEN FROM ITS DOSSIER

Every system separates the component that proposes from the component that decides. Every public repository of mine with code in it, on github.com/matrixbuilderops (2 written up, 3 small ones listed at the end). Descriptions are each repo's own words or verified write-ups; the Code line is what its biggest source files say they do and the grade its executed checks gave it; lines, tests and commits are measured by repo-dossier, with vendored, generated and duplicate code excluded.

AI EVALUATION, VERIFICATION & RED-TEAMING

groundwork — sitemap-mcp and filelens-mcp give agents structured awareness of sites and files; verifier-mcp, sorrell-engine-mcp and package-telemetry — the tool behind every download number on this résumé. 146 tests. TypeScript, Python
Python 5,735 · JavaScript 3,510 · TypeScript 2,958 lines · 146 tests passing (run) · 47 commits on 14 days (2026-08-13 → 2026-10-31) · NPM sitemap-mcp 1,130 downloads · NPM filelens-mcp 1,124 downloads · NPM verifier-mcp 230 downloads · NPM sorrell-engine-mcp 207 downloads · NPM package-telemetry 1,917 downloads · github.com/matrixbuilderops/groundwork
Code: index.ts: FileLens MCP Server Collapses multiple file-reading tool calls into single structured results · index.ts: SiteMap MCP Server Collapses website exploration into single structured results
Run: get · json · logs · check · profiles · blocks · unblock · never | Built on: @modelcontextprotocol/sdk, zod, websocket, webpipe, pytest, bs4, playwright, trafilatura, pydantic

OTHER WORK

workspace
Python 1,693 lines · 17 commits on 3 days (2026-09-09 → 2026-10-11) · github.com/matrixbuilderops/workspace
Code: nexus_os_corrected.py: The 'Multiplicative' Logic: In a traditional cluster, the 'Sum' is reduced by overhead (network latency, memory copying, context switching)
Run: --demo

Smaller repositories: cad-fea-pipeline, endless-watch, matrixbuilderops.

starforge — Token counter for AI-coding sessions, checked against an independently written counter by a different author: over 20,217 transcripts from 5 computers they agree exactly on every axis but output tokens, where they differ by 0.07%. 4,878 npm downloads (starforge-cli). JavaScript

SHIPPED TOOLING & PUBLIC ADOPTION (TELEMETRY VERIFIED)

- **46,179 lifetime downloads** across 11 published packages and 129 releases on PyPI and npm (live, 2026-11-01).
- **PyPI (33,873):** cli-enforcement (10.3k), cli-wikia (9.2k), cli-collective (7.4k), cli-fleet (7.0k).
- **npm (12,306):** starforge-cli (4.9k), starreckon (2.8k), package-telemetry (1.9k), sitemap-mcp (1.1k), filelens-mcp (1.1k), verifier-mcp (0.2k), sorrell-engine-mcp (0.2k).
- *Note:* Gross telemetry includes registry mirrors and CI runs; volunteered as directional proof of real-world usage, not audited human seats.

EDUCATION & CERTIFICATIONS

B.S., Integrative Studies — University of North Texas · 2019 · CompTIA Security+ · Network+ · A+

Alexander Sorrell — Verified Package Telemetry & Live Proof

Identity: Alexander Sorrell (@Alexander-Sorrell-IT) · alexander.sorrell.it@gmail.com · (817) 996-6123
Source: Real-time Registry API query via package_telemetry.py · Executed: 2026-11-01

LIVE TERMINAL TELEMETRY OUTPUT (PACKAGE_TELEMETRY.PY)

ALEXANDER SORRELL — PACKAGE METRICS & TELEMETRY

Type	Package	First Rel.	First Date	Current	Updated	Days Live	Last Day	Today	DL / Month	Lifetime DL
pip	cli-enforcement	v0.1.0	2026-06-29	v0.10.0	2026-10-11	125d	6	0	1,308	10,297
pip	cli-wikia	v0.10.0	2026-06-29	v0.21.0	2026-10-11	125d	6	0	1,247	9,187
pip	cli-collective	v0.1.0	2026-06-29	v0.7.0	2026-08-29	125d	2	0	719	7,411
pip	cli-fleet	v0.1.0	2026-06-29	v0.9.0	2026-10-11	125d	5	0	1,060	6,978
npx	starforge-cli	v0.1.0	2026-08-07	v0.14.0	2026-08-13	86d	5	0	67	4,878
npx	sitemap-mcp	v0.1.0	2026-08-13	v0.7.0	2026-08-28	80d	5	0	37	1,130
npx	filelens-mcp	v0.1.0	2026-08-13	v0.7.0	2026-08-28	80d	0	0	43	1,124
npx	starreckon	v0.14.0	2026-08-14	v0.31.0	2026-10-11	79d	15	0	718	2,820
npx	verifier-mcp	v1.0.0	2026-09-04	v1.0.0	2026-09-04	58d	2	0	31	230
npx	sorrell-engine-mcp	v1.0.0	2026-09-04	v1.0.0	2026-09-04	58d	1	0	22	207
npx	package-telemetry	v1.0.0	2026-09-06	v1.6.0	2026-10-11	56d	13	0	929	1,917

AVERAGES & SUMMARY

Category	Packages	Avg Days Live	Avg Last Day	Avg Today	Avg DL / Month	Avg Lifetime / Pkg	Total Lifetime DL
PyPI (pip)	4 pkgs	125.0d	4.8 / day	0.0 / day	1,083.5 / mo	8,468.2 / pkg	33,873
NPM (npx)	7 pkgs	71.0d	5.9 / day	0.0 / day	263.9 / mo	1,758.0 / pkg	12,306
ALL ECOSYSTEM	11 pkgs	90.6d	5.5 / day	0.0 / day	561.9 / mo	4,198.1 / pkg	46,179

PROGRAM DIRECTORY & WHAT THEY DO

Type	Package	What It Does / Architecture
pip	cli-enforcement	Model-agnostic hook-level behavioral enforcement engine for AI coding agents.
pip	cli-wikia	Offline, pip-installable reference wiki & search index for AI CLIs (Claude, Gemini, Copilot, DeepSeek).
pip	cli-collective	The Collective bundle: knowledge (cli-wikia) + control (cli-enforcement) + parallel execution (cli-fleet).
pip	cli-fleet	Hardware-aware parallel Claude Code agent orchestration via shared filesystem mailboxes.
npx	starforge-cli	Verification CLI auditing AI-coding telemetry with kernel-enforced no-egress proof.
npx	sitemap-mcp	MCP server providing website structural awareness without HTML token explosion.
npx	filelens-mcp	MCP server for single-call intelligent AST & structural file reading.
npx	starreckon	Local-only developer wrapped analyzer for disk logs (Claude Code, Cowork, Codex).
npx	verifier-mcp	Deterministic EU e-invoice compliance engine and paradox resolver (DevNetwork 2026).
npx	sorrell-engine-mcp	Browser-native neural OS & Tiny Model Swarm runtime (OpenAI WebMCP Challenge).
npx	package-telemetry	Cross-registry package telemetry engine for tracking PyPI and NPM package metrics, releases, and downloads.

Registry Ledger: packages_config.json • Canonical Output: package_telemetry_latest.json • Scope: PyPI & npm live registries

QR CODE TARGET ENGINE & REPRODUCIBLE VERIFICATION (NPX STARRECKON PROVE)

Routing Destination: github.com/Alexander-Sorrell-IT/starreckon · **Package:** npx starreckon

- **Reproducible Sandbox Execution:** npx starreckon prove prints the exact confinement command without running it, so you read the proof before you trust it. The command runs the scan under Apple sandbox-exec (Linux: unshare -rn) with a profile that denies network egress: (version 1) (allow default) (deny network*) (deny network-outbound) (deny network-inbound).
- **Kernel-Enforced Boundary:** The kernel refuses the socket (EPERM), so network egress is blocked below the application, not by an application-level boolean. A timeout is deliberately not counted as blocked.
- **The Central Engineering Thesis:** *The model proposes; deterministic executed checks own the verdict.*

PUBLIC INDEPENDENT VERIFICATION ENDPOINTS (ZERO INSTALL REQUIRED)

Anyone can query public package registries directly via cURL to independently verify these install metrics:

```
# 1. PyPI Verification (pypistats API, no key) — reproduces the PyPI figures above, which include mirrors
curl -s https://pypistats.org/api/packages/cli-enforcement/overall | jq '[.data[]|select(.category=="with_mirrors").downloads]|add'
#    without mirrors, the four PyPI packages total 9,892:
curl -s https://pypistats.org/api/packages/cli-enforcement/overall | jq '[.data[]|select(.category=="without_mirrors").downloads]|add'

# 2. npm Verification (via npm Registry API)
curl -s https://api.npmjs.org/downloads/point/last-month/starreckon | jq .downloads
curl -s https://api.npmjs.org/downloads/range/2020-01-01:2030-12-31/starforge-cli | jq '[.downloads[].downloads]|add'
```